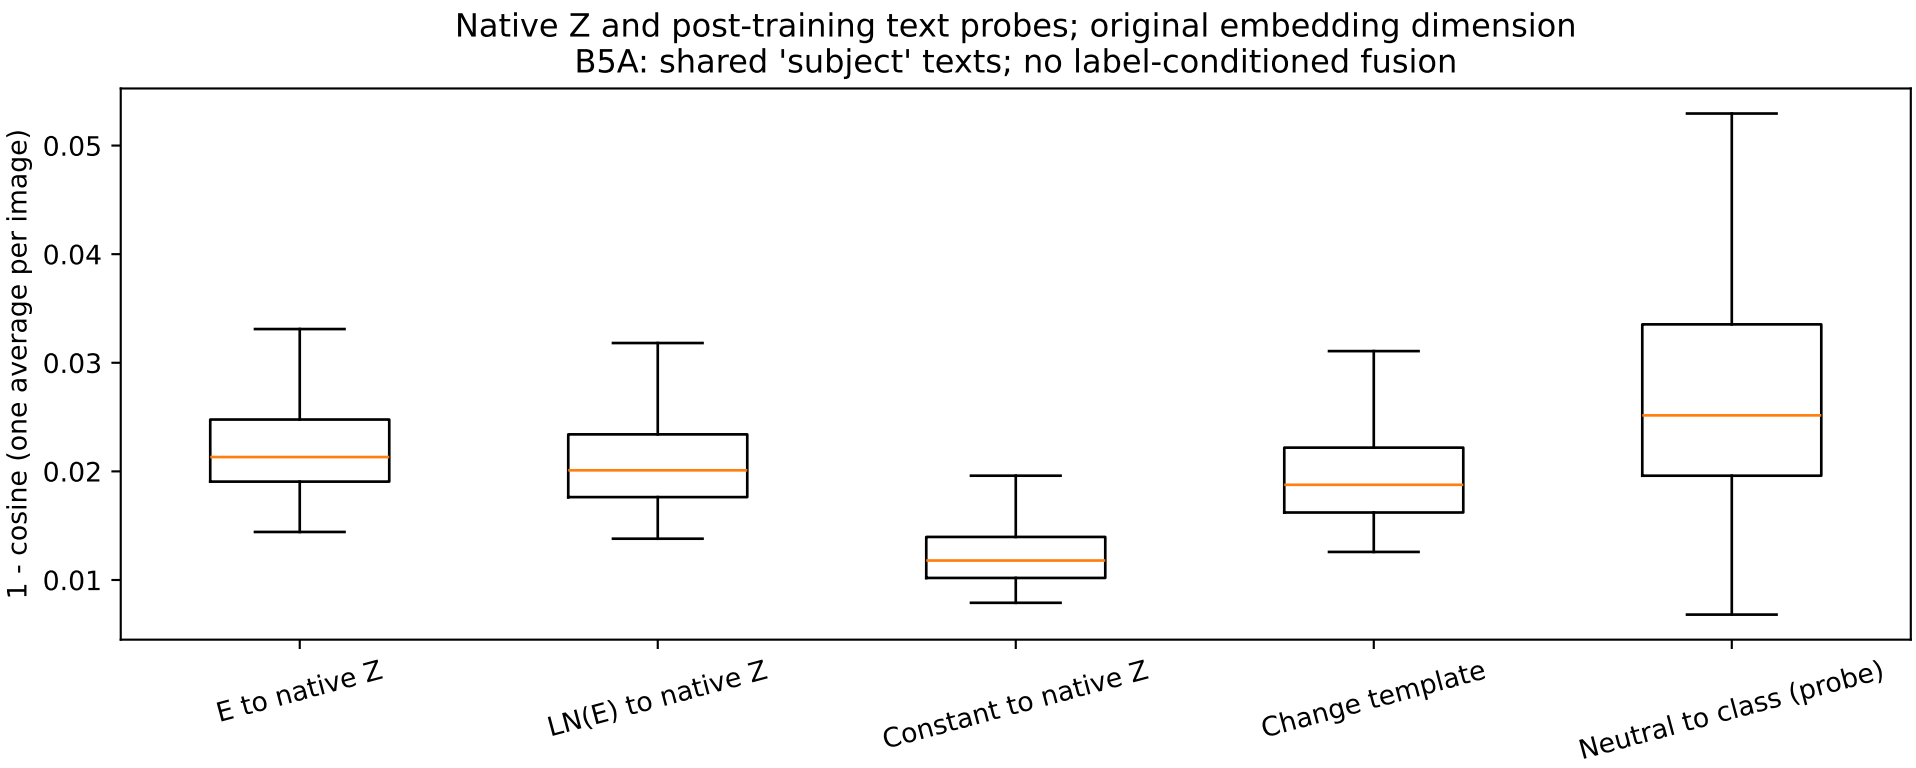
Native Z and post-training text probes; original embedding dimension
B5A: shared 'subject' texts; no label-conditioned fusion
1 - cosine (one average per image)
0.05
0.04
0.03
0.02
0.01
E to native Z
LN(E) to native Z
Constant to native Z
Change template
Neutral to class (probe)